

```
In [1]: import warnings
import itertools
import numpy as np
import matplotlib.pyplot as plt
warnings.filterwarnings("ignore")
plt.style.use('fivethirtyeight')
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_excel("Sample - Superstore.xls")
df
```

```
Out[2]:
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	...	Postal Code	Region	Product ID	Category	
0	1	CA-2016-152156	2016-11-08	2016-11-11	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	42420	South	FUR-BO-10001798	Furniture	E
1	2	CA-2016-152156	2016-11-08	2016-11-11	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	42420	South	FUR-CH-10000454	Furniture	
2	3	CA-2016-138688	2016-06-12	2016-06-16	Second Class	DV-13045	Darrin Van Huff	Corporate	United States	Los Angeles	...	90036	West	OFF-LA-10000240	Office Supplies	
3	4	US-2015-108966	2015-10-11	2015-10-18	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	...	33311	South	FUR-TA-10000577	Furniture	
4	5	US-2015-108966	2015-10-11	2015-10-18	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	...	33311	South	OFF-ST-10000760	Office Supplies	

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	...	Postal Code	Region	Product ID	Category	

9989	9990	CA-2014-110422	2014-01-21	2014-01-23	Second Class	TB-21400	Tom Boeckenhauer	Consumer	United States	Miami	...	33180	South	FUR-FU-10001889	Furniture	F
9990	9991	CA-2017-121258	2017-02-26	2017-03-03	Standard Class	DB-13060	Dave Brooks	Consumer	United States	Costa Mesa	...	92627	West	FUR-FU-10000747	Furniture	F
9991	9992	CA-2017-121258	2017-02-26	2017-03-03	Standard Class	DB-13060	Dave Brooks	Consumer	United States	Costa Mesa	...	92627	West	TEC-PH-10003645	Technology	
9992	9993	CA-2017-121258	2017-02-26	2017-03-03	Standard Class	DB-13060	Dave Brooks	Consumer	United States	Costa Mesa	...	92627	West	OFF-PA-10004041	Office Supplies	
9993	9994	CA-2017-119914	2017-05-04	2017-05-09	Second Class	CC-12220	Chris Cortes	Consumer	United States	Westminster	...	92683	West	OFF-AP-10002684	Office Supplies	A

9994 rows × 21 columns



```
In [3]: furniture = df.loc[df['Category'] == 'Furniture']
furniture['Order Date'].min(), furniture['Order Date'].max()
```

```
Out[3]: (Timestamp('2014-01-06 00:00:00'), Timestamp('2017-12-30 00:00:00'))
```

```
In [4]: cols = ['Row ID', 'Order ID', 'Ship Date', 'Ship Mode', 'Customer ID', 'Customer Name', 'Segment', 'Country', 'City',
furniture.drop(cols, axis=1, inplace=True)
furniture = furniture.sort_values('Order Date')
```

```
furniture.isnull().sum()
```

```
Out[4]: Order Date    0  
Sales          0  
dtype: int64
```

```
In [5]: furniture = furniture.groupby('Order Date')['Sales'].sum().reset_index()  
furniture
```

```
Out[5]:
```

	Order Date	Sales
0	2014-01-06	2573.8200
1	2014-01-07	76.7280
2	2014-01-10	51.9400
3	2014-01-11	9.9400
4	2014-01-13	879.9390
...
884	2017-12-24	1393.4940
885	2017-12-25	832.4540
886	2017-12-28	551.2568
887	2017-12-29	2330.7180
888	2017-12-30	323.1360

889 rows × 2 columns

```
In [6]: furniture = furniture.set_index('Order Date')  
furniture.index
```


```
Out[6]: DatetimeIndex(['2014-01-06', '2014-01-07', '2014-01-10', '2014-01-11',  
                        '2014-01-13', '2014-01-14', '2014-01-16', '2014-01-19',  
                        '2014-01-20', '2014-01-21',  
                        ...,  
                        '2017-12-18', '2017-12-19', '2017-12-21', '2017-12-22',  
                        '2017-12-23', '2017-12-24', '2017-12-25', '2017-12-28',
```

```
        '2017-12-29', '2017-12-30'],  
        dtype='datetime64[ns]', name='Order Date', length=889, freq=None)
```

```
In [8]: y_useful = furniture['Sales'].resample('MS').mean()  
        y_useful['2015':]
```


```
Out[8]: Order Date  
2015-01-01    978.328467  
2015-02-01    522.395667  
2015-03-01    781.236437  
2015-04-01    805.822962  
2015-05-01    624.996700  
2015-06-01    428.565500  
2015-07-01    719.706316  
2015-08-01    602.412012  
2015-09-01   1382.790684  
2015-10-01    632.980184  
2015-11-01   1286.701354  
2015-12-01   1049.355418  
2016-01-01    508.182867  
2016-02-01    356.868273  
2016-03-01    609.575810  
2016-04-01    695.373158  
2016-05-01    687.265227  
2016-06-01    816.910750  
2016-07-01    768.736412  
2016-08-01    734.307782  
2016-09-01   1135.953371  
2016-10-01    624.872474  
2016-11-01   1271.345152  
2016-12-01   1410.719808  
2017-01-01    397.602133  
2017-02-01    528.179800  
2017-03-01    544.672240  
2017-04-01    453.297905  
2017-05-01    678.302328  
2017-06-01    826.460291  
2017-07-01    562.524857  
2017-08-01    857.881889  
2017-09-01   1209.508583  
2017-10-01    875.362728  
2017-11-01   1277.817759  
2017-12-01   1256.298672  
Freq: MS, Name: Sales, dtype: float64
```


```
In [28]: y_useful.plot(figsize=(12, 6))
plt.show()
```



```
In [10]: decomposition = sm.tsa.seasonal_decompose(y_useful, model='additive')
trend = decomposition.trend
plt.figure(figsize=(12,16))
plt.subplot(411)
plt.plot(trend, label="Trend")
plt.legend()
Seasonal = decomposition.seasonal
plt.subplot(412)
plt.plot(trend, label="Seasonal")
plt.legend()
Residual = decomposition.resid
plt.subplot(413)
```

```
plt.plot(trend, label="Residual")  
plt.legend()  
plt.show()
```





```
In [11]: p = d = q = range(0, 2)
pdq = list(itertools.product(p, d, q))
```

```
In [12]: pdq
```

```
Out[12]: [(0, 0, 0),
(0, 0, 1),
(0, 1, 0),
(0, 1, 1),
(1, 0, 0),
(1, 0, 1),
(1, 1, 0),
(1, 1, 1)]
```

```
In [13]: seasonal_pdq = [(x[0], x[1], x[2], 12) for x in pdq]
```

```
In [14]: seasonal_pdq
```

```
Out[14]: [(0, 0, 0, 12),
(0, 0, 1, 12),
(0, 1, 0, 12),
(0, 1, 1, 12),
(1, 0, 0, 12),
(1, 0, 1, 12),
(1, 1, 0, 12),
(1, 1, 1, 12)]
```

```
In [15]: for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            model = sm.tsa.statespace.SARIMAX(y_useful, order=param, seasonal_order=param_seasonal, enforce_stationarity=False,
                                                enforce_invertibility=False)
            results = model.fit()
            print('ARIMA{0}{1} - AIC:{2}'.format(param, param_seasonal, results.aic))
```

```
except:
    continue
```

```
ARIMA(0, 0, 0)x(0, 0, 0, 12) - AIC:769.0817523205915
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
ARIMA(0, 0, 0)x(0, 0, 1, 12) - AIC:1354.0669954233208
```

```
ARIMA(0, 0, 0)x(0, 1, 0, 12) - AIC:477.7170130920899
```

```
ARIMA(0, 0, 0)x(0, 1, 1, 12) - AIC:302.27028997938197
```

```
ARIMA(0, 0, 0)x(1, 0, 0, 12) - AIC:497.2314433418337
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
ARIMA(0, 0, 0)x(1, 0, 1, 12) - AIC:1144.2225605403871
```

```
ARIMA(0, 0, 0)x(1, 1, 0, 12) - AIC:318.0047199116341
```

```
ARIMA(0, 0, 0)x(1, 1, 1, 12) - AIC:304.2488280301906
```

```
ARIMA(0, 0, 1)x(0, 0, 0, 12) - AIC:720.9252270758116
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
ARIMA(0, 0, 1)x(0, 0, 1, 12) - AIC:2695.913697357427
```

```
ARIMA(0, 0, 1)x(0, 1, 0, 12) - AIC:466.5607429809158
```

```
ARIMA(0, 0, 1)x(0, 1, 1, 12) - AIC:291.62613896732864
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
ARIMA(0, 0, 1)x(1, 0, 0, 12) - AIC:499.5869033885854
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
ARIMA(0, 0, 1)x(1, 0, 1, 12) - AIC:2347.5641565412016
```

```
ARIMA(0, 0, 1)x(1, 1, 0, 12) - AIC:319.98848769468657
```

```
ARIMA(0, 0, 1)x(1, 1, 1, 12) - AIC:291.8725576524215
```

```
ARIMA(0, 1, 0)x(0, 0, 0, 12) - AIC:677.894766843944
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
ARIMA(0, 1, 0)x(0, 0, 1, 12) - AIC:1320.66170690448
```

```
ARIMA(0, 1, 0)x(0, 1, 0, 12) - AIC:486.63785672282035
```

```
ARIMA(0, 1, 0)x(0, 1, 1, 12) - AIC:304.9671228167956
```

```
ARIMA(0, 1, 0)x(1, 0, 0, 12) - AIC:497.78896630044073
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood
```



```
d optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(0, 1, 0)x(1, 0, 1, 12) - AIC:1366.7797226279727
ARIMA(0, 1, 0)x(1, 1, 0, 12) - AIC:319.7714068109211
ARIMA(0, 1, 0)x(1, 1, 1, 12) - AIC:306.9113200151535
ARIMA(0, 1, 1)x(0, 0, 0, 12) - AIC:649.9056176817318
ARIMA(0, 1, 1)x(0, 0, 1, 12) - AIC:2508.436923977664
ARIMA(0, 1, 1)x(0, 1, 0, 12) - AIC:458.87055484828795
ARIMA(0, 1, 1)x(0, 1, 1, 12) - AIC:279.58062316811436
ARIMA(0, 1, 1)x(1, 0, 0, 12) - AIC:486.18329774426684
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(0, 1, 1)x(1, 0, 1, 12) - AIC:803.6652444374075
ARIMA(0, 1, 1)x(1, 1, 0, 12) - AIC:310.75743684175046
ARIMA(0, 1, 1)x(1, 1, 1, 12) - AIC:281.5576621461235
ARIMA(1, 0, 0)x(0, 0, 0, 12) - AIC:692.1645522067713
ARIMA(1, 0, 0)x(0, 0, 1, 12) - AIC:1350.6453425567156
ARIMA(1, 0, 0)x(0, 1, 0, 12) - AIC:479.46321478521355
ARIMA(1, 0, 0)x(0, 1, 1, 12) - AIC:304.2077675160951
ARIMA(1, 0, 0)x(1, 0, 0, 12) - AIC:480.92593679351927
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(1, 0, 0)x(1, 0, 1, 12) - AIC:1290.889297086618
ARIMA(1, 0, 0)x(1, 1, 0, 12) - AIC:304.46646750846253
ARIMA(1, 0, 0)x(1, 1, 1, 12) - AIC:304.5842692143838
ARIMA(1, 0, 1)x(0, 0, 0, 12) - AIC:665.7794442186472
ARIMA(1, 0, 1)x(0, 0, 1, 12) - AIC:2392.9009436960123
ARIMA(1, 0, 1)x(0, 1, 0, 12) - AIC:468.3685195815001
ARIMA(1, 0, 1)x(0, 1, 1, 12) - AIC:293.3422193965914
ARIMA(1, 0, 1)x(1, 0, 0, 12) - AIC:482.5763323877232
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(1, 0, 1)x(1, 0, 1, 12) - AIC:2436.877259630648
ARIMA(1, 0, 1)x(1, 1, 0, 12) - AIC:306.0156002132424
ARIMA(1, 0, 1)x(1, 1, 1, 12) - AIC:293.7513188135041
ARIMA(1, 1, 0)x(0, 0, 0, 12) - AIC:671.2513547541902
```

```
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(1, 1, 0)x(0, 0, 1, 12) - AIC:1180.8208704956457
ARIMA(1, 1, 0)x(0, 1, 0, 12) - AIC:479.2003422281134
```

```

ARIMA(1, 1, 0)x(0, 1, 1, 12) - AIC:300.21306116191005
ARIMA(1, 1, 0)x(1, 0, 0, 12) - AIC:475.34036587851494
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
  warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(1, 1, 0)x(1, 0, 1, 12) - AIC:1073.2557519174643
ARIMA(1, 1, 0)x(1, 1, 0, 12) - AIC:300.6270901345431
ARIMA(1, 1, 0)x(1, 1, 1, 12) - AIC:302.32649925046746
ARIMA(1, 1, 1)x(0, 0, 0, 12) - AIC:649.0318019835194
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
  warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(1, 1, 1)x(0, 0, 1, 12) - AIC:1370.2232859458645
ARIMA(1, 1, 1)x(0, 1, 0, 12) - AIC:460.4762687610177
ARIMA(1, 1, 1)x(0, 1, 1, 12) - AIC:281.3873006939415
ARIMA(1, 1, 1)x(1, 0, 0, 12) - AIC:469.5250354660838
C:\Users\U.R Computer\anaconda\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
  warnings.warn("Maximum Likelihood optimization failed to ")
ARIMA(1, 1, 1)x(1, 0, 1, 12) - AIC:1344.609794136266
ARIMA(1, 1, 1)x(1, 1, 0, 12) - AIC:297.7875439530794
ARIMA(1, 1, 1)x(1, 1, 1, 12) - AIC:283.36610143638575

```

In []:

In []:

```

In [26]: df_output = pd.Series(y_test[0:4],index=['Test Statistic','p-value','Lags Used','No. of Obs'])
         for i,j in y_test[4].items():
             df_output["Criticality is (%s)"%i] = j
         print(df_output)

```

```

Test Statistic      -5.191070
p-value             0.000009
Lags Used           10.000000
No. of Obs          37.000000
Criticality is (1%) -3.620918
Criticality is (5%) -2.943539
Criticality is (10%) -2.610400
dtype: float64

```

```

In [17]: from statsmodels.tsa.stattools import acf, pacf

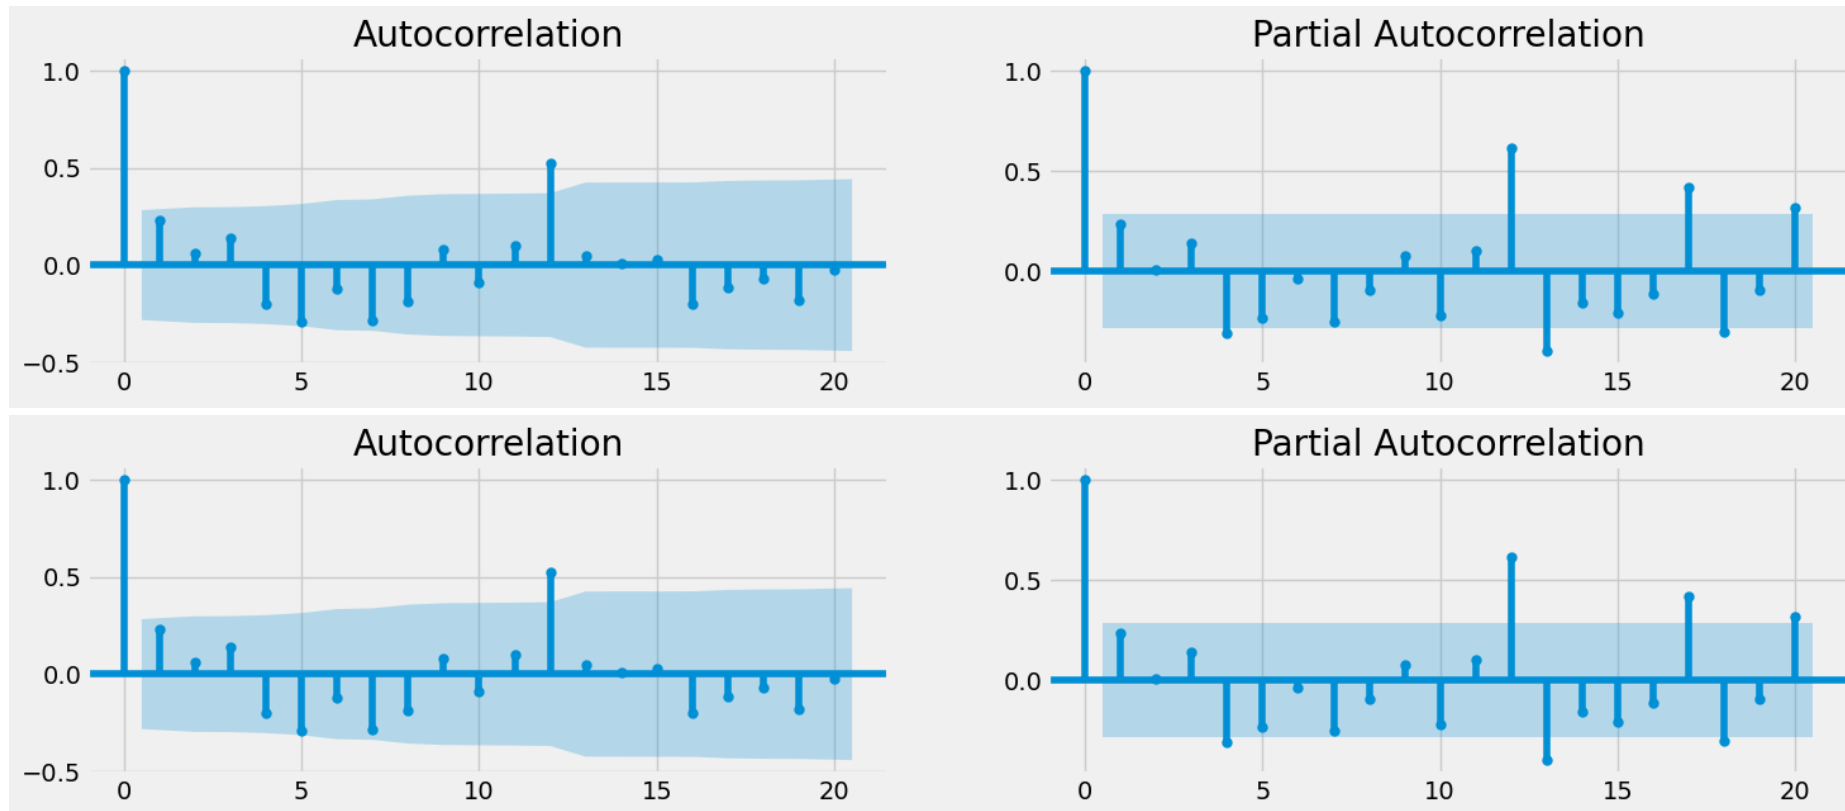
```

```

from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
# fig, (ax1,ax2) = plt.subplots(1,2,figsize=(10,5))
# fig = sm.graphics.tsa.plot_acf(y_useful.values.squeeze(),lags=20,ax=ax1)
# ax1.axhline(y=-1.96/np.sqrt(len(y_useful)))
# ax1.axhline(y=1.96/np.sqrt(len(y_useful)))
# fig = sm.graphics.tsa.plot_pacf(y_useful, lags=20, ax=ax2)
# ax2.axhline(y=-1.96/np.sqrt(len(y_useful)))
# ax2.axhline(y=1.96/np.sqrt(len(y_useful)))
fig, axes = plt.subplots(1,2,figsize=(16,3), dpi= 100)
plot_acf(y_useful, lags=20, ax=axes[0])
plot_pacf(y_useful, lags=20, ax=axes[1])

```

Out[17]:



```

In [18]: model = sm.tsa.ARIMA(y_useful,order=(2,0,2))
result = model.fit(dis=-1)
result

```

Out[18]: <statsmodels.tsa.arima_model.ARMAResultsWrapper at 0x21efbbf31f0>

```
In [19]: print(result.aic)
         print(result.bic)
```

693.5725330243508
704.7997390897982

```
In [20]: from sklearn import metrics
         results_pred = result.predict()
         metrics.mean_absolute_error(y_useful.values, results_pred.values)
```

Out[20]: 231.642903843485

```
In [52]: y_useful - results_pred
```

Out[52]:

Order Date	
2014-01-01	-305.286384
2014-02-01	-340.042123
2014-03-01	172.161118
2014-04-01	-236.917428
2014-05-01	-272.673920
2014-06-01	-18.867273
2014-07-01	-153.135170
2014-08-01	-271.649494
2014-09-01	288.759970
2014-10-01	-73.930893
2014-11-01	226.117066
2014-12-01	673.371605
2015-01-01	17.773495
2015-02-01	-301.594138
2015-03-01	35.468012
2015-04-01	11.963666
2015-05-01	-154.831394
2015-06-01	-318.918648
2015-07-01	24.100538
2015-08-01	-173.287443
2015-09-01	663.364133
2015-10-01	-323.984606
2015-11-01	587.776743
2015-12-01	73.201379
2016-01-01	-285.964330
2016-02-01	-402.065558

```

2016-03-01    -60.386323
2016-04-01    -59.109570
2016-05-01    -54.138546
2016-06-01     52.260155
2016-07-01    -19.138099
2016-08-01    -43.196669
2016-09-01    361.537749
2016-10-01   -253.599733
2016-11-01    547.332326
2016-12-01    463.489019
2017-01-01   -506.971833
2017-02-01   -173.358646
2017-03-01   -207.882565
2017-04-01   -251.289195
2017-05-01   -29.887843
2017-06-01    73.582003
2017-07-01   -225.027485
2017-08-01    136.902627
2017-09-01    391.926466
2017-10-01    -1.865678
2017-11-01    474.881569
2017-12-01    323.888756
Freq: MS, dtype: float64

```

```
In [53]: metrics.r2_score(y_useful, results_pred)
```

```
Out[53]: 0.06783298355618761
```

```
In [21]: result.summary()
```

```
Out[21]:
```


ARMA Model Results			
Dep. Variable:	Sales	No. Observations:	48
Model:	ARMA(2, 2)	Log Likelihood	-340.786
Method:	css-mle	S.D. of innovations	292.913
Date:	Wed, 14 Jul 2021	AIC	693.573
Time:	13:31:40	BIC	704.800
Sample:	01-01-2014	HQIC	697.815
	- 12-01-2017		

	coef	std err	z	P> z	[0.025	0.975]
const	785.4806	57.423	13.679	0.000	672.933	898.029
ar.L1.Sales	-0.0880	0.633	-0.139	0.889	-1.329	1.153
ar.L2.Sales	0.2451	0.311	0.787	0.431	-0.365	0.856
ma.L1.Sales	0.3534	0.618	0.572	0.568	-0.859	1.565
ma.L2.Sales	-0.1988	0.309	-0.643	0.520	-0.805	0.408


Roots				
	Real	Imaginary	Modulus	Frequency
AR.1	-1.8483	+0.0000j	1.8483	0.5000
AR.2	2.2072	+0.0000j	2.2072	0.0000
MA.1	-1.5236	+0.0000j	1.5236	0.5000
MA.2	3.3015	+0.0000j	3.3015	0.0000

```
In [22]: forc_arima = result.predict(start=36,end=96,dynamic=True)
plt.plot(forc_arima)
plt.plot(y_useful)
```

```
Out[22]: [<matplotlib.lines.Line2D at 0x21efd277280>]
```




```
In [47]: result.plot_predict(1,100)
x=result.forecast(steps=120)
```



```
In [56]: import statsmodels.api as sm
model=sm.tsa.statespace.SARIMAX(y_useful,order=(1, 1, 1),seasonal_order=(1,1,1,12))
results = model.fit()
forc_sarima = results.predict(start=36,end=96,dynamic=True)
```

```
plt.plot(forc_sarima,color="red")
plt.plot(y_useful,color="blue")
```

Out[56]: [<matplotlib.lines.Line2D at 0x21e80291190>]



In [24]: results.summary()

Out[24]:

SARIMAX Results						
Dep. Variable:		Sales		No. Observations:		48
Model:		SARIMAX(1, 1, 1)x(1, 1, 1, 12)			Log Likelihood	-238.291
Date:		Wed, 14 Jul 2021			AIC	486.582
Time:		13:31:46			BIC	494.359
Sample:		01-01-2014			HQIC	489.267
		- 12-01-2017				
Covariance Type:		opg				
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.0854	0.250	0.342	0.732	-0.404	0.575
ma.L1	-0.9984	13.948	-0.072	0.943	-28.337	26.340

ar.S.L12	0.0273	0.588	0.046	0.963	-1.124	1.179
ma.S.L12	-0.9921	63.664	-0.016	0.988	-125.772	123.788
sigma2	2.741e+04	1.84e+06	0.015	0.988	-3.58e+06	3.64e+06

Ljung-Box (L1) (Q):	0.05	Jarque-Bera (JB):	2.10
Prob(Q):	0.82	Prob(JB):	0.35
Heteroskedasticity (H):	0.54	Skew:	-0.37
Prob(H) (two-sided):	0.30	Kurtosis:	2.05

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [59]: `metrics.mean_absolute_error(y_useful.values, results.predict().values)`

Out[59]: 222.13466014555183

In []: